

INFFOREST Variable Importance on Random Forests

A Thesis
Presented to
The Division of Mathematics and Natural Sciences
Reed College

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts

Aurora Owens

May 2017

Approved for the Division
(Mathematics)

Andrew Bray

Acknowledgements

I want to thank a few people.

Preface

This is an example of a thesis setup to use the reed thesis document class.

Table of Contents

Chapter 1: Introduction	1
1.1 Trees and Random Forests	1
1.1.1 Trees	1
1.2 What We Mean When We Talk About Inference	2
1.2.1 Inferential vs Descriptive Statistics	2
1.3 Permutations and Populations	2
1.4 A Step Back	2
1.5 Inference on Random Forests	3
1.5.1 The Problem	3
1.5.2 Proposed solutions to this problem	3
Chapter 2: Simulations and Comparisons	5
2.1 Simulated Data	5
2.2 Models and Comparisons	7
Chapter 3: Random Forest Variable Importance	9
3.1 Breiman et al Introduce Permuted Variable Importance (1984)	9
3.1.1 Variable Importance on a Single Tree	9
3.1.2 Variable Importance for a Random Forest	10
3.2 Strobl et al Respond (2008)	10
Chapter 4: Chapter 4	13
4.1 INFFORESTS	13
4.2 Implementation In INFTREES and Results	14
4.2.1 Notes on the Implemetation	14
4.2.2 Results	16
Chapter 5: INFFOREST Comparisons With Other Methods	17

List of Tables

List of Figures

2.1	Density Graphs for V1 through V5 and a Plot of $Y \sim V4$, Correlation = .789	6
2.2	CART for the Model $Y \sim$, from D1	7
2.3	Simulated MSEtest Distributions of CART, Random, and Bagged Forests	8
4.1	Median Values of INFFOREST Variable Importance for mtry = 4,6 and 8	16
5.1	Median Values of INFFOREST, Conditionally Permuted, and Permuted Variable Importance	17

Abstract

The preface pretty much says it all.

Second paragraph of abstract starts here.

Dedication

You can have a dedication here if you wish.

Chapter 1

Introduction

1.1 Trees and Random Forests

1.1.1 Trees

Decision trees may be familiar to many with a background in social science as a convenient way to represent data and assist in decision making. Morgan and Sonquist (1963) derived a way for constructing trees motivated by the specific feature space of data collected from interviews and surveys. Unlike, agricultural data which involves mostly numeric variables, the data collected from interviews is mostly categorical. On top of this, the datasets Morgan and Sonquist dealt with had few participants, and a lot of data collected on each one. To add to their difficulties, there was reason to believe that there were lurking errors in the data that would be hard to identify and quantify. Lastly, many of the predictors were correlated. Morgan and Sonquist doubted that the additive assumptions of many models would be appropriate for this data. They noted that while many statistical methods would have difficulty accurately parsing this data, a clever researcher with quite a lot of time could create a suitable model simply by grouping values in the feature space and predicting that the response corresponding to these values would be an average of the observed responses given the grouped conditions. Their formalization of this procedure in terms of “decision rules” laid the ground work for future research on decision trees.

Later researchers proposed new methods for creating trees that improved upon the Morgan and Sonquist model. Leo Breiman et al (1984) proposed an algorithm called CART, *classification and regression trees*, that would allow trees to be fit on various types of data. An alternative to this method is conditional inference trees. Torsten Hothorn, Kurt Hornik, Achim Zeileis argue in their 2006 paper *Unbiased Recursive Partitioning: A Conditional Inference Framework*, CART has a selection bias toward variables with either missing values or a great number of possible splits. This bias can effect the interpretability of all tree models fit using this method. As an alternative to CART and other algorithms, Hothorn et al propose a new method, conditional inference trees.

There is a limit to the predictive capabilities of a single tree as they suffer from high variance. To alleviate this, aggregate methods called forests are often used

instead. They function by enlisting the help of many trees, and then by aggregating the responses over all of them. The two most common types of forests are bagged and random forests.

1.2 What We Mean When We Talk About Inference

1.2.1 Inferential vs Descriptive Statistics

A note should be made of the difference between inferential and descriptive statistics. This paper's aim is to describe a process of making inferential claims using random forests, not descriptive ones. Descriptive statistics describe the data at hand without making any reference to a larger data generating system that they come from. It follows that inferential statistics then make claims about the data generating system given the data.

1.3 Permutations and Populations

As stated in the introduction of the *Chronical of Permutations Statistical Methods* by KJ Berry et al, 2014, there are two models of statistical inference. One is the population model, where we assume that the data was randomly sampled from one (or more) populations. Under this model, we assume that the data generated follows some known distribution. "Under the population model, the level of statistical significance that results from applying a statistical test to the results of an experiment or a survey corresponds to the frequency with which the null hypothesis would be rejected in repeated random samplings from the same specified population(s)", (Berry et al, 2014).

The permutation family of methods, on the other hand, only assumes that the observed result was caused by experimental variability. The test statistics is first calculated for the observed data, then the data is permuted a number of times. The statistic is calculated after each permutation to derive a distribution of possible values. Then the original test statistic is tested against this distribution. If it is exceptionally rare, then there is evidence that our observation was not simply experimental variability.

1.4 A Step Back

A random forest R_f is the set of functions T_1, \dots, T_N where each T_j is a piece-wise function from the sample space Ω to the response space Φ . In general, Ω is defined by an $n \times p$ matrix where each column is a random variable and Φ is defined by an $n \times 1$ vector Y .

Each tree in a random forest, $T_j \in R_f$, is generated on a subset of both Ω and Φ called the training set. This training set is a bootstrapped sample of the original

dataset and is noted as B^t . It is then tested on a disjoint subset of Ω called the test set, \bar{B}^t , where $\bar{B}^t = \Omega \setminus B^t$. The image of T_j in Φ is called the predictions of T_j .

As outlined in the 1984 textbook, *Classification and Regression Trees*, Breiman, Friedman, Olshen, and Stone described their method for creating, pruning, and testing regression trees. There are essentially three steps: one, decide on a variable to split over, two, partition that variable space in two distinct partitions, and three, set our initial predictions for each partition to be mean value of the response according to the observed responses corresponding to the values in the partitions. Recursively, this process is repeated for each new partition until some stopping condition is reached. This is a top down, greedy algorithm that functions by creating as large a tree as possible and then is pruned down to prevent over fitting.

Random Forests are generated by fitting a large number of trees, each on a boosted sample of the data. The crucial difference, however, between the trees in CART and the trees in a random forest, is that at each node in a random forest, only a subset of the predictors are considered as candidates for possible splits. This decorrelates each tree from its neighbors, and decreases bias of the whole model while slightly increasing variance of each tree.

1.5 Inference on Random Forests

1.5.1 The Problem

Random forests create models with great predictive-, but poor inferential capabilities. After Morgan and Sonquist's initial development of decision trees, trees quickly moved to the domain of machine learning and away from statistics. Researchers focused on bettering predictions and improving run times and less on the statistics behind them. Inferential statistics with random forests is usually treated as a variable selection problem, and generally falls behind the predictions in importance. This has limited the applications of random forests in certain fields, as to many the question of "why" the data is the way it is, is just, if not more, important as the predictions. There are several means of performing descriptive statistics with random forests that could be interpreted incorrectly as attempting to answer this but without a statistically backed method for performing variable importance, the use of random forest is limited to prediction-only settings.

1.5.2 Proposed solutions to this problem

Breiman and Cutler proposed a method of permuted variable importance in their paper (cite) to answer this problem. Their method compares the variable importance for each variable in a tree-wise manner. For each tree, the permuted variable importance of the variable X_j is:

$$VI^t(x_j) = \frac{\sum_{i \in \bar{B}^t} (y - \hat{y})^2}{|\bar{B}^t|} - \frac{\sum_{i \in \bar{B}_p^t} (y - \hat{y}_p)^2}{|\bar{B}_p^t|}$$

Where \bar{B}^t is the out of bag sample for tree t , $|B|$ is the number of observations in that sample, \bar{B}_p^t is with X_j permuted, \hat{y} is the predicted outcome, and \hat{y}^t is the predicted outcomes after variable X_j has been permuted. This value is averaged over all the trees. It's important to note that if the variable X_j is not split on in the tree t , the tree-wise variable importance will be 0.

Creating a permutation-based method is certainly an attractive solution to our problem. One, it allows us to estimate the distribution of variable importance and generate a Z score under the null hypothesis that $PV = 0$.

$$VI_\alpha(x_j) = \frac{\sum_1^n treePV^t(x_j)}{\frac{\hat{\sigma}}{\sqrt{ntree}}}$$

Strobl et al from the University of Munich criticize this method in their 2008 technical report, *Danger: High Power! – Exploring the Statistical Properties of a Test for Random Forest Variable Importance*. One, this method has the downside of increasing power with increasing numbers of trees in the forest. This is a more or less arbitrary parameter which we would hope would not affect our importance estimates. Secondly, the null hypothesis under Breiman and Cutler's strategy is that the variable importance V for any variable X_j is not equal to zero given Y , the response. Because random forests are most often used in situations with multicollinearity that would make other methods like the linear model difficult, Strobl argues that any variable importance measure worth its salt should not be mislead by correlation within the predictors.

The researchers at the University of Munich published a fully fleshed response to the Breiman and Cutler method in 2008, titled *Conditional Variable Importance for Random Forests* that address these issues. Strobl et al propose restructuring the Breiman and Cutler algorithm to account for conditional dependence among the predictors. The null hypothesis is that $VI_\beta(X_j) = 0$ given the predictor Y and all other predictors X_1, \dots, X_n . This accounts for interactions between X_j and the other predictors.

This paper aims to provide a response to this method. The partitions are made from the random forest corresponding to the formula of Y X_1, \dots, X_n instead of a model of X_j X_1, \dots, X_n . This ignores the common situation where the predictors are correlated enough, they act as stand ins for each other, so that if one variable is heavily influential in a certain tree at predicting Y , the other variable will be forgotten all together.

Chapter 2

Simulations and Comparisons

2.1 Simulated Data

Tree-based methods shine in predictive situations with correlated predictors, although these situations can pose problems for inference. In a situation with correlated predictors X_1 and X_2 , and the treemodel we are considering is $Y \sim X_1 + X_2$, it is difficult to say how much of the modeled effect on Y is due to X_1 or X_2 . To illustrate this idea, compare a few existing methods, and explore methods of inference on tree based models three datasets will be simulated with different correlation structures. We will be focused more on the correlation structure between the predictors than on their relationships with the response and this will be reflected in the simulations.

To aid in comparisons between the methods, one of the simulated datasets considered in this paper will be generated from the same method as used in (Strobl et al, 2008b). Under this method, the 13 x 1000 data set, D_1 , has 12 predictors, V_1, \dots, V_{12} , where $V_j \sim N(0, 1)$. The first four are, however, block correlated to each other with $\rho = .9$. They are related to Y by the linear equation:

$$Y = 5 \cdot V_1 + 5 \cdot V_2 + 2 \cdot V_3 + 0 \cdot V_4 + -5 \cdot V_5 + -5 \cdot V_6 + 0 \cdot V_7 + 0 \cdot \dots + E, E \sim N(0, \frac{1}{2})$$

Note that the coefficients for V_7, \dots, V_{12} are all zero.

Table 1: Correlation of V_1, \dots, V_7 and Y

	V1	V2	V3	V4	V5	V6	V7	y	beta
V1	1.000	0.915	0.908	0.907	-0.034	0.006	0.012	0.829	5
V2	0.915	1.000	0.914	0.914	-0.020	-0.001	-0.001	0.830	5
V3	0.908	0.914	1.000	0.903	-0.017	-0.007	0.007	0.808	2
V4	0.907	0.914	0.903	1.000	-0.002	-0.015	0.023	0.789	0
V5	-0.034	-0.020	-0.017	-0.002	1.000	0.044	0.005	-0.388	-5
V6	0.006	-0.001	-0.007	-0.015	0.044	1.000	-0.005	-0.364	-5
V7	0.012	-0.001	0.007	0.023	0.005	-0.005	1.000	-0.141	-2

As can be seen from the last column in the table, “beta”, although V_4 was not included in the model $Y \sim V_1, \dots, V_{12}$, it has a strong correlation with more influential

predictors V_1, \dots, V_3 insures that it still shows a strong, empirical linear correlation with Y . A linear model would likely *overstate* the effect of V_4 on Y .^{1 2}


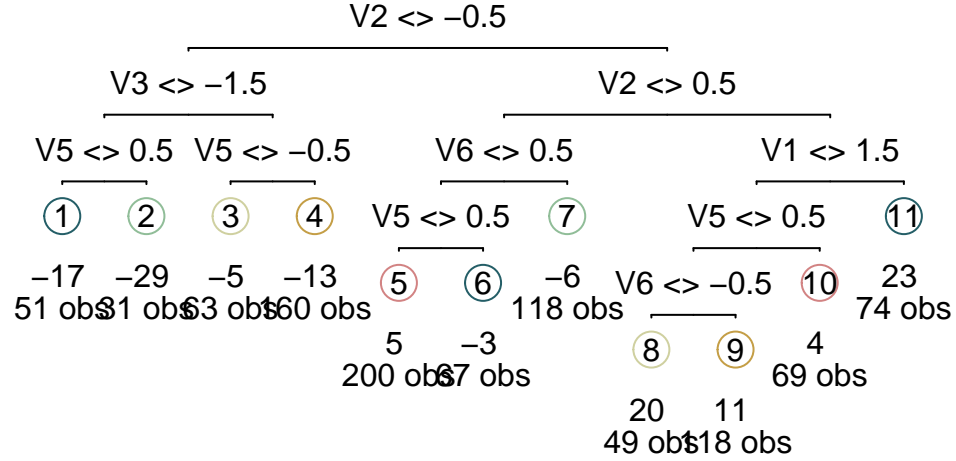


Figure 2.1: Density Graphs for V_1 through V_5 and a Plot of $Y \sim V_4$, Correlation = .789

As can be seen in Figure 1 the densities of V_1, \dots, V_5 are all very similar due to the way they were generated.

Figure 2.2: CART for the Model $Y \sim$, from D_1

2.2 Models and Comparisons

CART: Regression Trees

A single CART tree representing the model $Y \sim X_1, \dots, X_{12}$ is easy enough to understand. Starting at the very top of the tree, predictions can be made based on the values of the leaves (or ending nodes) given the requirements of the path to get there. Trees can be quite variable, so to get a better idea of the differences between the methods let's run a simulation.

Algorithm 1 Simulation Scheme 2.1

- 1: **for** $i \leq 1000$ **do**
 - 2: Randomly sample $\frac{2}{3}$ of the observations in D_1 to a training set, $D_{1,train}^i$. The other observations, $x \in D_1, x \notin D_{1,train}^i$ form the testing set $D_{1,test}^i$
 - 3: Fit a tree, T^i , to the data under the model $Y \sim X_1, \dots, X_2$ using the observations in D_1^i
 - 4: Calculate the MSE_{test} of the model using the equation: $MSE_{test} = \frac{1}{n} \sum (y_j - \hat{y}_j)^2$
 - 5: **end for**
-

Note that n is the number of observations in $D_{1,test}^i$, $y_j \in D_{1,test}^i$, $\hat{y}_j \in T^i(D_{1,test}^i)$

¹A brief note on uncertainty is needed here. It's true that in this setting we can say that V_4 is actually unimportant to understanding Y , but in situations with real data this is profoundly more difficult to parse. Often like in the social science situations that Morgan and Sonquist encountered, the real relationship between correlated predictors is complicated and often there is some theoretical backing or other insight that is gained to include variables that may not be important to the model.

²Another point that could be said is that, no V_4 is not unimportant, V_1, V_2 , and V_3 are just stand ins for the real star, V_4 , as they are nearly the same ($\rho \sim 1$). Then the real relationship represented here is $Y \sim (5 + 5 + 2) \cdot V_4 + -5 \cdot V_5 + -5 \cdot V_6 + -2 \cdot V_7$. This model is not unsuccessful in capturing the structure of the data, and this is typically the practice used to model data with highly correlated predictors. If this seems philosophically satisfying to you, the rest of this thesis may seem a bit inconsequential.

for $1 \leq j \leq n$ This produces one distribution of MSE_{test} for CART.




Figure 2.3: Simulated MSE_{test} Distributions of CART, Random, and Bagged Forests

The distribution of 100 CART trees' MSE_{test} in the above figure is roughly normal with a variance of `var(testmseC)`. There is a fair amount of variability in a single tree, they are heavily dependent on fluctuations in the starting data set. As mention briefly in the introduction, bagged forests present one solution to this problem. To create a bagged forest, as outlined in *An Introduction to Statistical Learning* by James, Witten, Hastie and Tibshirani, 2013, many bootstrapped samples are taken from the initial dataset and trees are fitted to them. The final predictions are, then, averaged over all of the trees. This ensures that while each tree has high variance, when they are aggregated the variance will decrease. As one can see, the values of MSE_{test} for the bagged forest were entirely below the MSE_{test} for the trees and the variance was much smaller. As random forests are unbiased, they can be much smaller than their bagged forest cousins without sacrificing accuracy.

Chapter 3

Random Forest Variable Importance

3.1 Breiman et al Introduce Permuted Variable Importance (1984)

3.1.1 Variable Importance on a Single Tree

Breiman et al in *Classification and Regression Trees* (1984) propose a method for variable importance for individual trees that stems from their definition of \tilde{s} , a surrogate split. Surrogate splits help Breiman et al deal with several common problems: missing data, masking, and variable importance. They are defined using logic that resembles that behind random forests.

Definitions

Assume the standard structure for tree models. Let D be the dataset composed of $D = Y, X_1, \dots, X_p$, where the model we would like to estimate is of the form $T : Y \sim X_1, \dots, X_p$. For any node $t \in T(D)$, s^* is the best split of the node into daughters t_r and t_l . Take $X_i \in D$ and let S_i be the set of all of the splits on X_i in T . Then set \bar{S}_i equal to the complement of S_i , $\bar{S}_i = S_i^c$. For any possible split $s_i \in S_i \cup \bar{S}_i$, s_i will split the node t into two daughters, $t_{i,l}$ and $t_{i,r}$. Count the number of times that s^* and s_i , while splitting differently, generate the same left daughter t_l as N_{LL} and the number of times they generate the same right daughter as N_{RR} . Then the probability that a case falls within $t_L \cap t'_L$ is $P(t_L \cap t'_L) = \sum_j \frac{\pi(j)N_j(LL)}{N_j}$ and the probability that a case falls within $t_R \cap t'_R$ is $P(t_R \cap t'_R) = \sum_j \frac{\pi(j)N_j(RR)}{N_j}$. Where $\pi(j)$ is the prior assumed for the j th variable. Finally, the probability that a surrogate split predicts s^* is $P(s^*, s_M) = (t_R \cap t'_R) + P(t_L \cap t'_L)$. Then the surrogate split is the value of s^* that maximizes this probability. It is denoted \tilde{s}

A surrogate split \tilde{s} , is one that estimates the best possible univariate split s^* on node t .

Defintion: Variable Importance, Single Tree

$$VI_{tree}(X_i, T) = \sum_{t \in T} \Delta RSS(\tilde{s}_i, t)$$

Or the decrease of RSS attributable to X_i across the tree T . In *Classification and Regression Trees*, Breiman et al, outline several potential problems with this method that they do not attempt to solve. First, that this is only one of a number of reasonable ways to define variable importance. Second, the variable importances for variables X_1, \dots, X_p can be effected by outliers or random fluctuations within the data. (Ch 5.3)

3.1.2 Variable Importance for a Random Forest

One way to define variable importance for a random forest follows directly from Breiman et al's definition for a single tree. Recall that each tree in a random forest is fit to a bootstrapped sample of the original observations. To estimate the test error, therefore, no cross validation is needed - each tree is simply tested against the test set of observations that were not in that tree's initial training set. To determine variable importance for a predictor X_j , we look at the RSS of the each tree's prediction that did not split on X_j . These values are then averaged over the subset forest that did not include X_j . A large value would imply that in trees that included X_j , the predictive capabilities were increased.

To formalize that idea, let's refer to the set of trees that did not consider X_j , $T_{x_j}^c$. Now, $T_{x_j}^c \subset R$, the random forest. The subset of the original data that will be tested on each tree, t , is \bar{B}^t . The dimensions of \bar{B}^t are $\nu_t \times p$, where p is the number of predictors and $\nu \leq n$. The number of trees in $T_{x_j}^c$ is μ where $\mu \leq ntree$

Now, base variable importance is:

$$VI_\alpha(X_j, R) = \sum_{t \in T_{x_j}^c} \frac{1}{\nu_t} RSS(t, \bar{B}_t)$$

However, this method poses some problems. Namely, while variable importance for random forests is more stable than for the variable importance values for CART, (this is because the model is less variable in general), it is lacking the traditional inferential capabilities of other regression models. In an effort to derive a p-value for variable importance values, Breiman 2001b, describes a *permuted variable importance* or VI_β that does not utilize $T_{x_j}^c$.

Again, a large variable importance value suggests that X_j is a valuable predictor for the model.

3.2 Strobl et al Respond (2008)

Strobl et al (2008) respond to Breiman's method with one main argument: the null hypothesis implied by the permutation distribution utilized in permuted variable importance is that X_i is independent of Y **and** $X_j \notin X_1, \dots, X_p$ so the null hypothesis will be rejected in the case where X_j is independent of Y but not some subset of the

Algorithm 2 Permuted Variable Importance for Random Forests, VI_β

Fit a random forest, R on the dataset D fitting the model $Y \sim X_1, \dots, X_p$.
for each $X_i \in X_1, \dots, X_p$ **do**
 for each $t \in R$ **do**
 Calculate: $\Phi_o = \frac{1}{\nu_t} RSS(t, \bar{B}^t)$
 Permute X_i . Now find $\Phi^* = \frac{1}{\nu_t} RSS(t, \bar{B}_t^*)$
 The difference between these values, $\Phi^* - \Phi_o$, is the variable importance for X_j on t ,
 end for
 Average over all $t \in R$

$$VI_\beta(X_j) = \frac{1}{ntree} \sum^{ntree} \Phi^* - \Phi_o$$

$$VI_\beta(X_j) = \frac{1}{ntree} \sum^{ntree} \frac{1}{\nu_t} RSS(t, \bar{B}_t^*) - \frac{1}{\nu_t} RSS(t, \bar{B}^t)$$

end for

other predictors. As correlation among the predictors is very common in data sets that are used for random forests, this is a large problem for Breiman's method.

To alleviate this difficulty, Strobl et al propose a permutation scheme under the null hypothesis that X_j given it's relationship with the other predictors is independent of Y .

There are several ways mentioned in ref:@strobl2008 to choose the set of predictors Z to condition X_j upon. Z might be chosen due to outside theory about the problem or Z might include every $p - 1$ predictor possible. In that paper's simulations section as well as in this one's, Z is chosen as the set of predictors with empirical correlation $\geq .2$

Algorithm 3 Conditional Variable Importance for Random Forests, VI_γ

- 1: Fit a random forest, R on the dataset D fitting the model $Y \sim X_1, \dots, X_p$.
- 2: **for** each $t \in R$ **do**
- 3: Calculate: $\Psi_o = \frac{1}{\nu_t} RSS(t, \bar{B}^t)$
- 4: **for** each $X_i \in X_1, \dots, X_p$ **do**
- 5: Select $Z \in X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_p$ to condition on when permuting X_j
- 6: Use the cutpoints on each variable in Z to create a grid on X_j
- 7: Permute X_j with respect to this grid
- 8: Now find $\Psi^* = \frac{1}{\nu_t} RSS(t, \bar{B}_t^*)$
- 9: The difference between these values, $\Psi^* - \Psi_o$, is the variable importance
- 10: **end for**
- 11: Average over all $t \in R$

$$VI_\gamma(X_i, R) = \frac{1}{ntree} \sum \Psi^* - \Psi_o$$

$$VI_\gamma(X_i, R) = \frac{1}{ntree} \sum \frac{1}{\nu_t} RSS(t, \bar{B}_t^*) - \frac{1}{\nu_t} RSS(t, \bar{B}^t)$$

12: **end for**

Chapter 4

Chapter 4

4.1 INFFORESTS

The INFFOREST variable importance is a method of permuted variable importance not unlike that of conditionally permuted variable importance (Algorithm 3). Permuted variable importance is calculated at the tree level, using the partitions on X_j from a tree created to predict the model $X_j \sim X_1, \dots, X_{j-1}, X_{j+1}, \dots, X_p$. This auxiliary tree is fit by considering all $p - 1$ predictors at each split and so may be quite large or quite small depending on the richness of the correlation structure around X_j . The auxiliary tree is also fit using the OOB sample for the tree at question. If the auxiliary tree results in a single leaf, i.e. there are no splits, then X_j is permuted blindly, without partitions. If the auxiliary tree results in two leaves, there will be two partitions on X_j to permute X_j within, and so on. After permuting X_j within these partitions, the RSS is calculated for that tree using the OOB sample of predictors, including the now-permuted X_j . The absolute difference of the RSS after permutation and the RSS with the untouched OOB sample is INFFOREST variable importance for that tree. Note that for this reason, the INFFOREST variable importance is always greater than or equal to zero, and is standardized by the max INFFOREST variable importance value given by that tree. As the variable importance values are calculated for each tree for each variable, once the method is completed there is a distribution of potential variable importance values for X_j , one for each tree. These distributions may or may not be normal, depending on the multicollinearity of the predictors. The INFFOREST variable importance algorithm works as follows:

INFFOREST variable importance operates under the null hypothesis that Y is independent of X_j given the correlation structure of X_j and the other $p - 1$ predictors, or that the true INFFOREST variable importance for X_j is 0. The alternative hypothesis is that Y and X_j are not independent given the correlation structure of X_j and the other predictors or that the INFFOREST variable importance for X_j is greater than zero. After INFFOREST values have been computed for the entire forest, they are treated as samples from the population of possible INFFOREST values for X_j given the random forest R_f , and a significance test can be under the null hypothesis.

Algorithm 4 INFForests, $VI_{inf}(R)$

-
- 1: Fit a random forest, R on the dataset D fitting the model $Y \sim X_1, \dots, X_p$.
 - 2: **for** each $X_i \in X_1, \dots, X_p$ **do**
 - 3: **for** each $t \in R$ **do**
 - 4: Calculate: $\Xi_o = \frac{1}{\nu_t} RSS(t, \bar{B}^t)$
 - 5: Calculate a tree T_i that predicts $X_i \sim X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_p$ using the subset of the observations used to fit t
 - 6: Permute the subset of X_i contained in \bar{B}_t with respect to the set of partitions P_{xi} from T_i .
 - 7: Now find $\Xi^* = \frac{1}{\nu_t} RSS(t, \bar{B}_t^*)$
 - 8: The difference between these values, $\Xi^* - \Xi_o$, is the variable importance for X_i on t
 - 9: **end for**
 - 10: Test the null hypothesis that 0 is the likely value of $\frac{1}{\nu_t} RSS(t, \bar{B}_t^*)$ using the distribution of values of Ξ^* gathered from each tree in R
 - 11: **end for**
-

4.2 Implementation In INFTREES and Results

4.2.1 Notes on the Implemetation

Implementing the INFFOREST and therefor the INFTREES algorithms, required creating a suite of functions to create trees and random forests. The trees are fit following the standard two-part CART-like algorithm.¹ The function chooses a variable to split on with linear correlation with respect to Y , but instead of looking for correlations above a certain threshold which is common, it chooses the variable with the highest correlation when compared to its peers. This alleviates the situation where a variable with a non-linear relationship would be passed over again and again. The splitting is done via minimization of the following function with respect to i :

$$RSS_{node}(i, X, Y) = RSS_{leaf}(Y|X < i) + RSS_{leaf}(Y|X \geq i)$$

$$RSS_{leaf} = \sum (y - \hat{y})^2$$

$$\hat{Y} : \hat{y} \in \hat{Y} : \hat{y} = E(Y), \text{ where } |\hat{Y}| = |Y|$$

This function considers the regression case only, and only numeric predictors. Leafs are created when the resultant split would be unsatisfactory, i.e. at least one daughter node would have five members or less. This generates very large trees - a quality that is not an issue in random forests but may be problematic in a stand-alone setting. At this time, there is also no function to prune the trees.

¹A brief note on uncertainty is needed here. It's true that in this setting we can say that V_4 is actually unimportant to understanding Y , but in situations with real data this is profoundly more difficult to parse. Often like in the social science situations that Morgan and Sonquist encountered, the real relationship between correlated predictors is complicated and often there is some theoretical backing or other insight that is gained to include variables that may not be important to the model.

Table 2: A Home-Grown Tree on $Y = X_1 + X_2 + X_3 + X_4$

var	n	dev	ypred	split.cutleft
X2	50	5958.56398616138	-3.32099458459633	1.89262782336418
leaf	14	1683.79172385909	15.2166924465285	0
X1	36	1771.9972397028	-10.5300950967004	-0.758420438327835
X1	27	1061.15524312941	-5.71617332730625	-0.251307488862014
leaf	17	720.083908357861	-2.51280971153014	0
leaf	10	341.071334771552	-11.1618914741256	0
leaf	9	239.837381494473	-24.971860404883	0

The tree output is read in the following way: each row corresponds to a node of the tree which considers **n** observations. The mean of the Y values included in the node are **ypred**. If there is an optimal and allowable split,² then the chosen variable, **var**, and the RSS_{node} , **dev**, are recorded.³ The value of the variable in question that acts as the split point is recorded as **split.cutleft**. If there is no split on the node in question, then **var** will be recorded as **<leaf>** and the **dev** value will be the value of RSS_{leaf} at this node.

The tree output is read roughly from top to bottom, with a coda in the middle. The first row corresponds to the first node, or the node that includes the entire dataset. The second row is the beginning of the right subtree or the right daughter of first node. This pattern continues, favoring the right daughter, until a leaf is reached. The left daughter of the first node is found after all of the splits off of the right daughter have finished but is easily identified as the row with a value of **n** that is exactly the difference between the **n** values of the first two rows. In the case where the right daughter contained many more observations of the original dataset, there may be a node within the right subtree that contains the same number of observations as the left daughter of the first node. In this case, the left daughter is simply the second row with this property. The pattern of following the right daughter until a leaf is reached continues with the left subtree.

The *INFTREE* function follows the algorithm referenced earlier. The partitions on X_j are generated by fitting a tree, T , to the model $X_j \sim X_1, \dots, X_{j-1}, X_{j+1}, \dots, X_p$ and calculating the predictions $T(X_1, \dots, X_{j-1}, X_{j+1}, \dots, X_p)$. Then permuting X_j with respect to the partitions on X_j given by those predictions. For example, if $x_j \in X_j$ and the value of $T(x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_p)$ corresponding to x_j is α , x_j is permuted along with the other values of X_j that also have $T(x_1, \dots, x_{j-1}, x_{j+1}, \dots, x_p)$ corresponding to α .

The values of $INFFOREST(X_j)$ are scaled in the following way: since the *INFFOREST* function computes the *INFTREES*, (or the difference in post and pre

²Another point that could be said is that, no V_4 is not unimportant, V_1, V_2 , and V_3 are just stand ins for the real star, V_4 , as they are nearly the same ($\rho \sim 1$). Then the real relationship represented here is $Y \sim (5 + 5 + 2) \cdot V_4 + -5 \cdot V_5 + -5 \cdot V_6 + -2 \cdot V_7$. This model is not unsuccessful in capturing the structure of the data, and this is typically the practice used to model data with highly correlated predictors. If this seems philosophically satisfying to you, the rest of this thesis may seem a bit inconsequential.

³Recall that we only allow splits to take place that split the data into two groups, each with more than five members.

permutation RSS), values in a tree-wise manner, each tree's values are divided by the maximum value. This ensures that the values are between zero and one, and that in each tree one variable is clearly deemed the *most important*.

4.2.2 Results




Figure 4.1: Median Values of INFFOREST Variable Importance for $mtry = 4, 6$ and 8

As in (ref Strobl et al 2008), the median INFFOREST variable importance scores are reported here for the dataset D_1 .⁴

As noted in several publications (Strobl et al, Breiman, Intro to Stat Learning), random forests structure is dependent on the value of $mtry$.⁵ INFFOREST variable importance remains fairly consistent as $mtry$ fluctuates.

⁴It's the convention to call the RSS_{node} the deviance at a node N , but, of course, this only makes sense when the node is a leaf.

⁵A great deal of effort was undertaken by the author to find the definitive, authentic CART algorithm. This implementation follows the rough strokes set out in the 1984 text *Classification and Regression Trees* to the best of the author's ability and may not be exactly the algorithm found in R packages like 'tree()'

Chapter 5

INFFOREST Comparisons With Other Methods




Figure 5.1: Median Values of INFFOREST, Conditionally Permuted, and Permuted Variable Importance

INFFOREST holds its own amongst the other methods described in Chapter 4. The conditional permuted variable importance, when ran on the same random forest, had more difficulty parsing out the situation with paired variables than INFFOREST, only selecting the first two variables to be important. Permuted variable importance

does not pick up on the correlation structure within the predictors and deems $V2$ and $V6$ unimportant.

INNFOREST and conditional permuted variable importance both ignored the unfluent predictors that were not correlated with $V1, V2, V3$ and $V4$. This is a situation that also occurred in the simulation run in Strobl et al (2008b).